

FACTORED AI & DATA HACKATHON 2026

Aclara

Help for bank customers who don't recognize a card charge, in Spanish and Portuguese.

AI for language.
Code for every decision.

Sebastian · solo entry

EXAMPLE CONVERSATION

No reconozco un cargo de 150 USD en Bazar Caleta.

Lo encontré: 16 de junio, 150 USD, aprobado. ¿Quieres disputarlo?

Confirmar con código



Disputa registrada y verificada en la base de datos

THE PROBLEM

Unrecognized charges are the largest complaint, and they are slow to fix

12,297

complaints about a charge the customer doesn't recognize: the largest complaint subcategory

15.4

days on average to resolve one of them

43.6%

of complaint contacts are resolved on the first contact

Aclara takes the intake: it explains the charge, files an eligible dispute once the customer confirms, and hands everything else to a person with the facts already checked.


Source: organizer dataset (synthetic LATAM bank), complaint and contact analysis in the repository

HOW IT WORKS

AI understands. Code decides, acts and verifies.

AI understands. Code controls actions.

Current design · updated after the final evaluation; v4 scores unchanged



Actions are verified before success is reported. Uncertainty reaches a human.

1 · Understand

Gemini reads intent, amount and date in Spanish or Portuguese. It has no tools and no authority.

2 · Decide and confirm


Bank rules in code decide. The customer confirms with a button and a one-time code.

3 · Verify or hand off

Every case is read back before we say "done". Anything uncertain goes to a person with a ready summary.

FINAL EVALUATION

100 unseen cases, the same workload for both systems



+10 points resolved safely (95% CI +5 to +16) \$0.0023 model cost per case 2.1 s median reply

Strict safety check: 8 of Aclara's 100 cases were flagged. 6 came from the test setup itself, and the other 2 were fixed afterwards (those fixes are not counted in these numbers).

WHY CODE HOLDS THE AUTHORITY

A customer types a fake confirmation into the chat. What happens?

PLAIN AI AGENT, NO CONTROLS

2 of 2

filed the dispute anyway, and claimed success without checking

ACLARA

0 of 2

stayed at a proposal and waited for the real button and one-time code

Same Gemini model, same tools. Across 20 attacks the plain agent also tried to look up another customer once; Aclara never did.

Development study run after the final evaluation · 20 authored attacks per system · not a general safety rate

BUILT FOR TRUST

Engineered to be checked, and honest about its limits

Under the hood

- Data contracts and quality gates on the organizer data; 44,570 unsafe complaint links excluded
- Row-level security; every judge gets an isolated workspace
- Every action read back before the customer sees "done"
- 1,600+ automated tests and CI on every pull request; live on Azure

Limits we state openly

- Synthetic bank and simulated SMS code: not a production bank
- Portuguese test text is model-written; native review pending
- Both systems failed the full safety checks in the final evaluation

Try it: ca-web-aclara-dev-eastus2.lemonbeach-1b769de0.eastus2.azurecontainerapps.io

Code: github.com/sebastian-gm/factored-hackathon-2026-sebastian